

Freight Rate Prediction - Assessment Report

Candidate submission for Spotter

August 2026

1. Summary

The task is to predict freight posted_rate. I built two models - a **full-feature** model for the 12,000 graded validation loads and a **reduced** model for the fixed December chart - validated them with a **time-based** split, and submitted an **equal-weight blend** of the three strongest models.

Headline validation numbers (out-of-fold, time-based CV):

Metric	Value
RMSE	629.3
MAE	114.3
MAPE	5.1%
R ²	0.826

2. Data exploration - key findings

- **48,000** training loads spanning **Jan-Oct 2025**; the validation set is **Nov-Dec 2025**. The problem is therefore a *forecast*: predict later months from earlier ones.
- **64** pickup cities, **64** delivery cities, **4,014** unique lanes, and **3** equipment types (Dry Van 57%, Reefer 25%, Flatbed 18%).
- posted_rate ranges **\$57-\$25,533** and is strongly driven by distance and market_index.
- **The pricing is close to linear**. A plain Ridge regression matched or beat gradient boosting at baseline - a strong signal that rate is roughly distance × per-mile price, without severe nonlinearity.

3. Data-quality issues and how I handled them

- **Synthetic coordinates**. Each city has exactly one lat/long (no dupes), but the coordinates are geographically wrong (e.g. “Los Angeles” is mislocated). I computed the straight-line (haversine) distance and found it correlates with the provided distance at **r = 0.9995** (ratio ≈ 1.18, a plausible road detour factor). Conclusion: coordinates are internally consistent but **redundant with distance**, so I did not rely on them (kept only a detour-ratio feature).
- **Weight sign errors**. There are **292 negative weights**. Their magnitudes match ordinary weights, so I repaired them with `abs(weight)` and retained negative/missing audit indicators as model features.
- **Detached label errors**. A robust equipment + log-distance rate-per-mile baseline exposes **677 labels (1.41%)** separated from the main residual distribution by a completely empty band. These

rows are excluded from model fitting only. Every held-out row remains in the reported metrics, so the headline score still represents the scorer's all-row evaluation policy.

- **Irreducible noise floor.** Different model families now converge near **RMSE $\approx$ 629-632** ($R^2 \approx 0.83$). When independent models hit the same wall, that wall is noise the features cannot explain. I therefore focused gains on *proportional* error (MAPE) rather than chasing RMSE.

4. Validation approach and data split

I used a **time-based (expanding-window) split**, never a random one - a random split would leak future information into training and overstate accuracy.

Each fold trains on all earlier months and validates on a single later month, holding out **August, then September, then October 2025**. This mirrors the real task (Nov-Dec is unseen future) and gives an honest estimate.

All **target encoding is fitted inside each fold** (train rows only) and applied to the held-out rows, so no target information leaks across the split. The label-quality screen is also derived from each training fold only. It never filters or alters the validation target.

Two separate models are trained because the inputs differ:

Model	Used for	Columns available
Full	12,000 graded loads	all (incl. market_index, quote_signal, coords)
Reduced	December chart	only the 6 shared columns

5. Feature engineering

- **Calendar:** month, weekday, day-of-month, week-of-year, weekend & US-holiday flags, and smooth seasonal sin/cos of day-of-year (so a future month stays in-range for tree models).
- **Numeric:** distance, log-distance, repaired weight, negative/missing-weight indicators, weight-per-mile, and interactions (distance $\times$ equipment, and for the full model distance $\times$ market_index, distance $\times$ quote_signal).
- **Categorical:** equipment one-hot; **smoothed target encoding** for lane, pickup and delivery (leak-safe, fitted within folds).
- **Target framing:** the single biggest win was predicting **log(rate)** - it cut MAPE from **8.4% to 5.1%** while RMSE stayed near the noise floor.

6. Model bake-off

Roster: Ridge, RandomForest, XGBoost, LightGBM, CatBoost, MLP, and TabPFN (TabPFN run on Kaggle - see note below). Light Optuna tuning was used on the strongest candidates. The final comparison below uses the same saved hyperparameters after the fold-safe data corrections, so the improvement is attributable to the corrected data path rather than a new search budget.

Candidate	RMSE	MAE	MAPE
LightGBM	629.9	119.1	5.37%
CatBoost	630.1	117.1	5.31%
Ridge	631.6	124.5	5.62%
XGBoost	632.3	122.5	5.51%
Three-model blend	629.3	114.3	5.13%

TabPFN (Kaggle), for completeness. The assessment machine has no GPU, so TabPFN ran on Kaggle. Two real constraints applied: the current package gates its newest weights behind a license token that a non-interactive kernel cannot accept, and Kaggle's GPU/torch build raised a CUDA kernel-image mismatch - so it ran on **CPU with a 3,000-row training context**. Under that constraint TabPFN was **not competitive**: CV RMSE $\approx 1,141$, MAPE $\approx 37\%$. This is expected - TabPFN excels on *small* datasets, whereas this one is large and near-linear, a regime gradient boosting and linear models capture fully using the complete training set. It is therefore reported but excluded from the final blend.

7. Final model

Per the plan, I compared the best single model against an equal-weight blend of the top 3 on out-of-fold predictions. The **blend (LightGBM + CatBoost + Ridge)** won and was refit on the screened training rows to produce `validation_predictions.csv` (12,000 rows, mean $\approx \$2,375$).

8. Reproducibility

uv-managed, Python 3.11 pinned. Run `uv sync` then `uv run python -m src.finalize` to regenerate predictions and the chart. See `README.md` for the full command list and `SPEC.md` for design decisions. Four focused tests cover weight repair, label screening, training-only row selection, and the fixed December output schema. The supplied scorer validates all 12,000 submission rows and all 31 December rows.

9. December prediction

December is a *future* month, and its file has only 6 columns, so I used the **reduced CatBoost (log)** model with calendar features. Because the model keys off weekday and seasonality - not a raw time index - the daily line moves realistically instead of flatlining.




Figure 1: December 2025 predicted rate (Lexington → Fort Wayne).

The clear **weekly cycle** (mid-week peaks, weekend dips, **\$807-\$836**) is the calendar features at work.